

# Approximate Policy Iteration

Reinforcement Learning

Roberto Capobianco



SAPIENZA  
UNIVERSITÀ DI ROMA

# Recap

# Finite-Horizon MDPs

---

Slightly different formulation:

$$(S, A, R, T, \mathbf{H}, \mu_0)$$


(time-horizon)  $\mathbf{H} \geq 0$  and  $s_0 \sim \mu_0$  (initial state distribution)

We consider time-dependent policies  $\pi$

$$\pi = \{\pi_0, \pi_1, \pi_2 \dots \pi_{H-1}\}$$

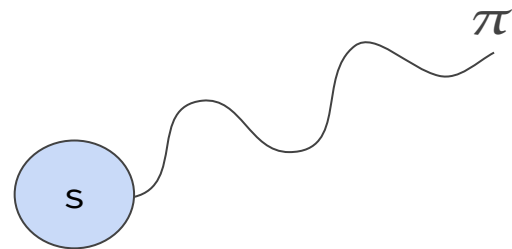
Actions might be different for the same state depending on  $t$

Very common in control!



# Finite-Horizon MDP: V & Q

---

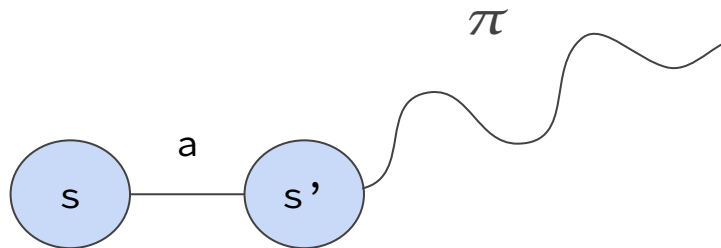


$$V_h^\pi(s) = \mathbb{E}_\pi[\sum_{\square=h}^{H-1} r(s_\square, a_\square)]$$

where  $s_h=s$ ,  $a_\square=\pi_\square(s_\square)$  and  $s_{\square+1}\sim P(\cdot|s_\square, a_\square)$

$$Q_h^\pi(s, a) = \mathbb{E}_\pi[\sum_{\square=h}^{H-1} r(s_\square, a_\square)]$$

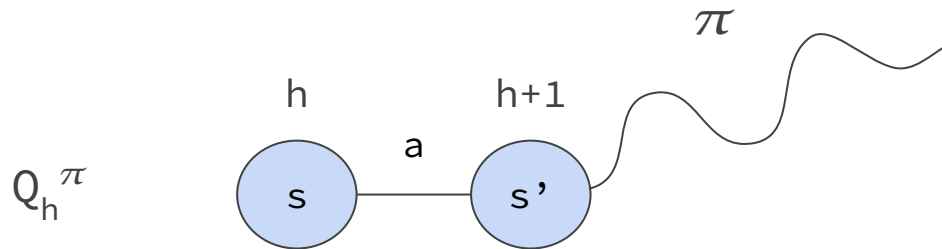
where  $s_h=s$ ,  $a_h=a$ ,  $a_\square=\pi_\square(s_\square)$  and  $s_{\square+1}\sim P(\cdot|s_\square, a_\square)$



# Finite-Horizon MDP: Bellman Equation

— — —

$$Q_h^\pi(s, a) = r(s, a) + \mathbb{E}_{s' \sim p(\cdot | s, a)} [V_{h+1}^\pi(s')]$$



# Finding the Optimal Policy

---

$$\pi^* = \{\pi_0^*, \pi_1^*, \pi_2^* \dots \pi_{H-1}^*\}$$

Let's reason backwards in time and apply dynamic programming:

$$Q_{H-1}^*(s, a) = r(s, a)$$

$$\pi_{H-1}^*(s) = \operatorname{argmax}_a Q_{H-1}^*(s, a)$$

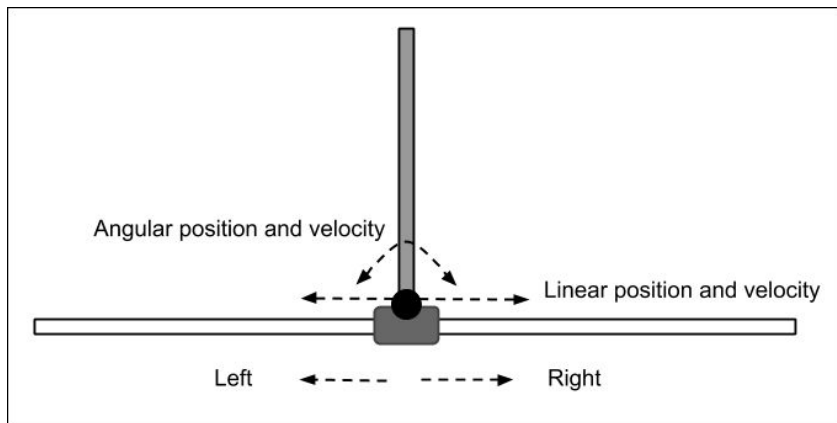
$$V_{H-1}^*(s) = \max_a Q_{H-1}^*(s, a) = Q_{H-1}^*(s, \pi_{H-1}^*(s))$$



# Control Problems

---

So far, we assumed discrete state and action spaces, but what about cartpole:



- **state:** angular pos & vel, linear pos & vel
- **action/control:** force applied on the cart
- **goal:** find the control policy which minimizes the long term cost  $c$



# Optimal Control

---

Given a dynamical system with a non-linear transition function  $f$ , state  $x$  in  $\mathbb{R}^d$  and control  $u$  in  $\mathbb{R}^k$ , we want to find a control policy  $\pi$  such that

$$\text{minimize } \mathbb{E}_{\pi}[c_H(x_H) + \sum_{h=0}^{H-1} c_h(x_h, u_h)]$$

$$\text{where } u_h = \pi(x_h) \text{ and } x_0 \sim \mu_0$$

Now this seems very familiar! Can we treat it as a  
Finite-Horizon MDP





# Bellman's Curse of Dimensionality

---

- $n$ -dimensional (discrete) state space
- The number of states grows exponentially in  $n$

In practice discretization is useful, but it is only computationally feasible up to 5 or 6 dimensional state spaces

Let's try to work directly in continuous space, starting from simplified problems



# Linear Systems

---

Consider a system of this kind:

$$x_{t+1} = Ax_t + Bu_t$$

This is our  
transition function!

- $x_t$  state at time  $t$
- $u_t$  control (i.e., action) at time  $t$

$A$  in  $\mathbb{R}^{d \times d}$ ,  $B$  in  $\mathbb{R}^{d \times k}$



# Quadratic Cost Function

---

Consider a cost function of this kind

$$c(x_t, u_t) = x_t^T Q x_t + u_t^T R u_t$$

alternative notation  
 $g(x_t, u_t)$

- $Q$  in  $\mathbb{R}^{d \times d}$  and  $R$  in  $\mathbb{R}^{k \times k}$  square matrices
- $Q$  and  $R$  positive definite

As a result, there is a non-zero cost for any non-zero state with all-zero control



# LQR algorithm

---

- Initialize  $P_H$  (at 0 or  $Q$ )
- Starting from  $h = H-1$ , backwards
  - Set  $K_h^* = -(R+B^T P_{h+1} B)^{-1} B^T P_{h+1} A$
  - Compute  $u = \pi_h^*(x_t) = K_h^* x_t$
  - Set  $P_h = (Q + K_h^{*T} R K_h^* + (A + B K_h^*)^T P_{h+1} (A + B K_h^*))$
  - Set  $J_h^* = x_t^T P_h x_t$

Riccati  
Equation

This is the Value Iteration update done in closed form: it is always the same and solves this particular continuous-state system with a quadratic cost



# LQR Extensions

---

Extensions to the LQR make it more generally applicable to:

- Affine systems
- Systems with stochasticity
- Regulation around non-zero fixed point for non-linear systems
- Trajectory following for non-linear systems
- ...



# End Recap



SAPIENZA  
UNIVERSITÀ DI ROMA

# Policy Iteration

— — —

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$
- Different from Value Iteration that was outputting values

Procedure:

1. Start with a random guess  $\pi_0$  (can be deterministic or stochastic)
2. For  $t=0, \dots, T$ :

$Q^\pi(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^\pi(s')]$

  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  for all  $s, a$
  - b. Do **policy improvement** as  $\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$  for all  $s$

This algorithm only makes progress, and the performance progress of the policy is monotonic



# State Visitation Probability

— — —

What's the probability of visiting state  $s$ ,  $a$  at time  $t$  according to  $\pi$  starting at  $s_0$ ?

$$d_{s_0}^{\pi}(s, a) = (1-\gamma) \sum_{h=0}^{\infty} \gamma^h \mathbb{P}^{\pi}_h(s, a; s_0)$$

$$\mathbb{P}^{\pi}_t(s, a; s_0) = \sum_{a_0, s_1, a_1, \dots, s_{t-1}, a_{t-1}} \mathbb{P}^{\pi}(s_0, a_0, \dots, s_t=s, a_t=a)$$

$$\mathbb{P}^{\pi}(s_0, a_0, \dots, s_t, a_t) = \pi(a_0 | s_0) p(s_1 | s_0, a_0) \pi(a_1 | s_1) p(s_2 | s_1, a_1) \dots p(s_t | s_{t-1}, a_{t-1}) \pi(a_t | s_t)$$





# State Visitation Probability

---

What's the probability of visiting state  $s$ ,  $a$  at time  $t$  according to  $\pi$  starting at  $s_0$ ?

$$d_{s_0}^{\pi}(s, a) = (1-\gamma) \sum_{h=0}^{\infty} \gamma^h \mathbb{P}^{\pi}_h(s, a; s_0)$$

Note that  $d_{s_0}^{\pi}$  is an infinite mixture

$$\mathbb{P}_t^{\pi}(s, a; s_0) = \sum_{a_0, s_1, a_1, \dots, s_{t-1}, a_{t-1}} \mathbb{P}^{\pi}(s_0, a_0, \dots, s_t=s, a_t=a)$$

$$\mathbb{P}^{\pi}(s_0, a_0, \dots, s_t, a_t) = \pi(a_0 | s_0) p(s_1 | s_0, a_0) \pi(a_1 | s_1) p(s_2 | s_1, a_1) \dots p(s_t | s_{t-1}, a_{t-1}) \pi(a_t | s_t)$$



# Approximate Policy Iteration

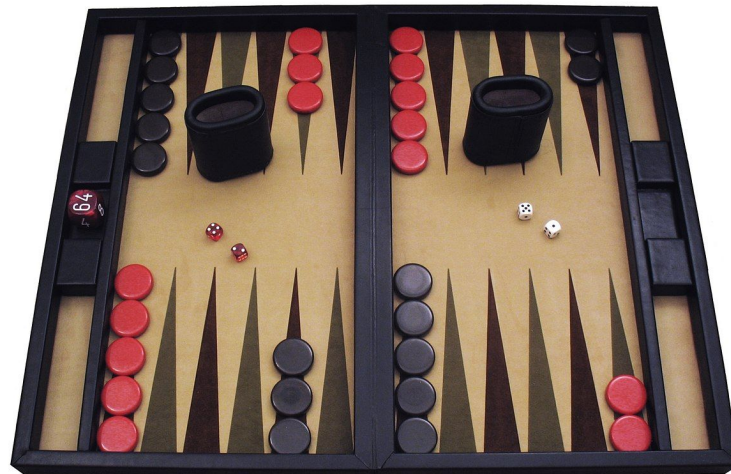
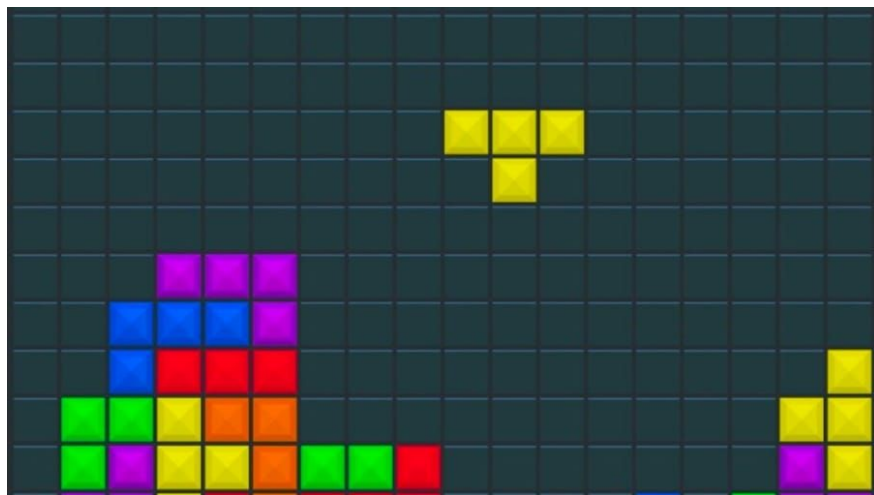
---

What if the state-space is large and we cannot do exact or iterative policy evaluation for all states?

# Approximate Policy Iteration

---

What if the state-space is large and we cannot do exact or iterative policy evaluation for all states?



# Approximate Policy Iteration

---

What if the state-space is large **or continuous** and we cannot do exact or iterative policy evaluation for all states?



# Approximate Policy Iteration

---

What if the state-space is large and we cannot do exact or iterative policy evaluation for all states?

**Assumptions:** the (infinite-horizon) MDP is still known, but the state-space is too large to just enumerate all states and compute  $V^\pi(s)$



# Approximate Policy Iteration

---

What if the state-space is large and we cannot do exact or iterative policy evaluation for all states?

## Assumptions:

$(S, A, R, T, \gamma, \mu_\theta)$  is given

$Q$  is in  $[0, 1/(1-\gamma)]$



# Approximate Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$

Procedure:

1. Start with a random guess  $\pi_0$
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi_t}$  ~~for all  $s, a$~~

$$Q^{\pi}(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^{\pi}(s')]$$



# Approximate Policy Iteration

---

- Outputs policies at every iteration:  $\{\pi_0, \pi_1, \pi_2 \dots \pi_T\}$

Procedure:

1. Start with a random guess  $\pi_0$
2. For  $t=0, \dots, T$ :
  - a. Do **policy evaluation** and compute  $Q^{\pi^t}$  ~~for all  $s, a$~~ 
$$Q^{\pi}(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot | s, a)} [V^{\pi}(s')]$$
  - b. Do **policy improvement** as  $\pi_{t+1} = \operatorname{argmax}_a Q^{\pi^t}(s, a)$  ~~for all  $s$~~   
~~argmax is still doable, we can still enumerate actions or discretize them~~





# Approximate Policy Evaluation

---

We build an **approximation**  $V^\pi$  of the true value function  $V^\pi$

If the approximation is close to the true value, then the optimal policy will be close-to-optimal



# Approximate Policy Evaluation

---

We build an **approximation**  $V^\pi$  of the true value function  $V^\pi$

If the approximation is close to the true value, then the optimal policy will be close-to-optimal

Error bounds exist, but we will skip them for your happiness :)



# Approximate Policy Evaluation

---

We build an **approximation**  $V^{\pi}$  of the true value function  $V^{\pi}$

If the approximation is close to the true value, then the optimal policy will be close-to-optimal

**Approximation for large state-spaces is needed to generalize among states and avoid looking at the whole  $S$**



# Approximate Policy Evaluation

---

We build an **approximation**  $V^{\pi}$  of the true value function  $V^{\pi}$

If the approximation is close to the true value, then the optimal policy will be close-to-optimal

**Approximation for large state-spaces is needed to generalize among states and avoid looking at the whole  $S$**

We use a function approximator



# Approximate Policy Evaluation

---

We build an **approximation**  $V^\pi$  of the true value function  $V^\pi$

If the approximation is close to the true value, then the optimal policy will be close-to-optimal

**Approximation for large state-spaces is needed to generalize among states and avoid looking at the whole  $S$**

We use a function approximator

e.g., linear approximators, neural nets, non-parametric, etc.



# Approximate Policy Evaluation

— — —

To be fair, we can directly approximate  $Q$ , so let's do that



# Approximate Policy Evaluation

— — —

To be fair, we can directly approximate  $Q$ , so let's do that

**Note that this also means that we can also get rid of the assumption of knowing the MDP**



# Approximate Policy Evaluation

— — —

To be fair, we can directly approximate  $Q$ , so let's do that

What do we need?





# Data and Least Square Regression

---

To be fair, we can directly approximate  $Q$ , so let's do that

What do we need?

**DATA  $D = \{s_i, a_i, y_i\}_{i=1}^N$  with  $y$  being our label!**

with those we can then use least-square regression to extract a function  $Q$  in the family of functions

$$S \times A \rightarrow [0, 1/(1-\gamma)]$$



# Data and Least Square Regression

---

To be fair, we can directly approximate  $Q$ , so let's do that

What do we need?

**DATA  $D = \{s_i, a_i, y_i\}_{i=1}^N$  with  $y$  being our label!**

with those we can then use least-square regression to extract a function  $Q$  in the family of functions

$$Q: S \times A \rightarrow [0, 1/(1-\gamma)]$$

$$\operatorname{argmin}_{Q \in \mathcal{Q}} \sum_{i=1}^N (Q(s_i, a_i) - y_i)^2$$



# Data and Least Square Regression

---

$$\operatorname{argmin}_{\mathbf{Q}} \sum_{i=1}^N (\mathbf{Q}(\mathbf{s}_i, \mathbf{a}_i) - y_i)^2$$

This is just a regression problem, which

- is numerically tractable
- has generalization bounds



# Supervised Learning Digression



# Supervised Learning



## Regression

What is the temperature going to be tomorrow?

PREDICTION

84°

Fahrenheit  
°F



## Classification



Will it be Cold or Hot tomorrow?

PREDICTION


COLD

HOT

Fahrenheit  
°F



Classification



Regression



# Supervised Learning: Regression

---

Given a ***data distribution*** from which we sample points  $x_i$  and labels  $y_i = f(x_i) + \epsilon_i$ , with  $\mathbb{E}[\epsilon_i] = 0$  and  $|\epsilon_i| \leq c$ , we want to approximate  $f$  using a finite set of data (dataset):

$$f^\wedge = \operatorname{argmin}_{f^\wedge \in F} \sum_{i=1}^N (f^\wedge(x_i) - y_i)^2$$

$$\text{with } F = \{f^\wedge: X \rightarrow \mathbb{R}\}$$



# Supervised Learning: Regression

---

Given a **data distribution**  $D$  from which we sample points  $x_i$  and labels  $y_i = f(x_i) + \epsilon_i$ , with  $\mathbb{E}[\epsilon_i] = 0$  and  $|\epsilon_i| \leq c$ , we want to approximate  $f$  using a finite set of data (dataset):

Empirical  
Risk  
Minimizer

$$f^\wedge = \operatorname{argmin}_{f^\wedge \in F} \sum_{i=1}^N (f^\wedge(x_i) - y_i)^2$$

with  $F = \{f^\wedge: X \rightarrow \mathbb{R}\}$

We can generalize under the same data distribution

$$\mathbb{E}_{x \sim D} (f^\wedge(x) - f(x))^2 \leq \delta \text{ with } \delta \text{ small}$$



# Supervised Learning: Distribution Mismatch

— — —

$\mathbb{E}_{x \sim D'} (f^{\wedge}(x) - f(x))^2$  can be huge!

If  $D' \neq D$






# End

## Supervised Learning Digression



# Oscillation from Distribution Change

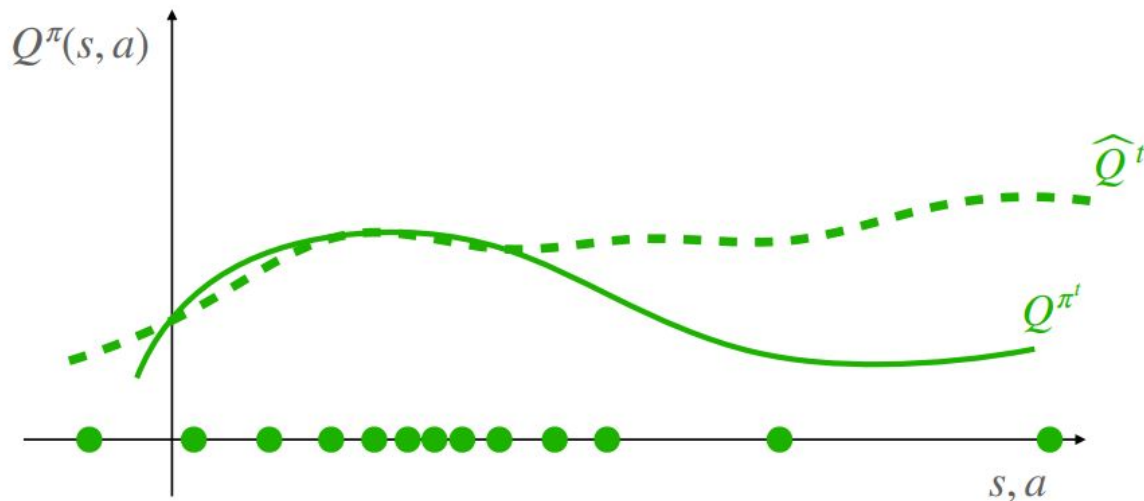


Credits: Wen Sun

Green dots:  $(s, a)$  from  $\pi^t$   
Red dots:  $(s, a)$  from  $\pi^{t+1}$



# Oscillation from Distribution Change




Credits: Wen Sun

Green dots:  $(s, a)$  from  $\pi^t$   
Red dots:  $(s, a)$  from  $\pi^{t+1}$



# Oscillation from Distribution Change



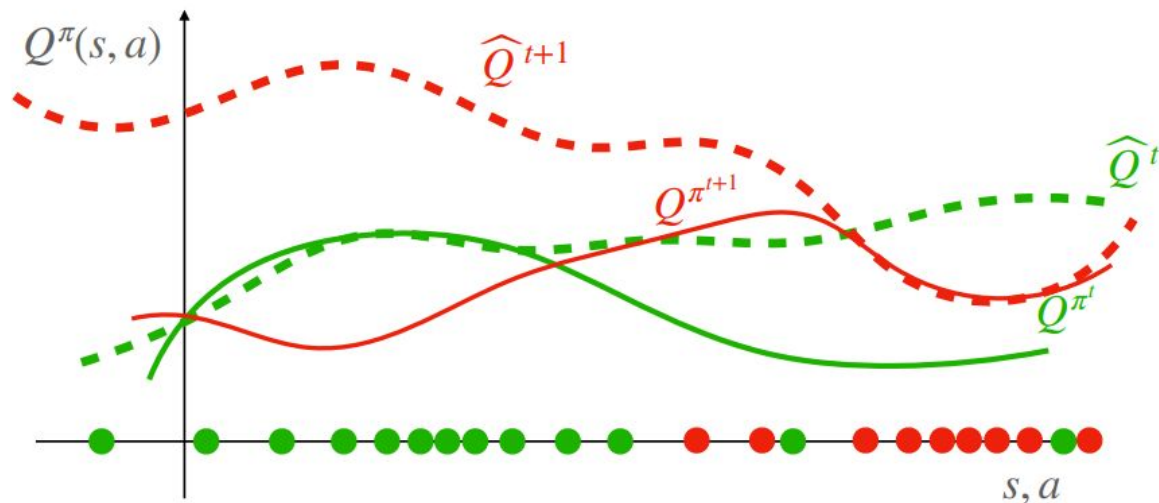
Credits: Wen Sun

Green dots:  $(s, a)$  from  $\pi^t$   
Red dots:  $(s, a)$  from  $\pi^{t+1}$



SAPIENZA  
UNIVERSITÀ DI ROMA

# Oscillation from Distribution Change

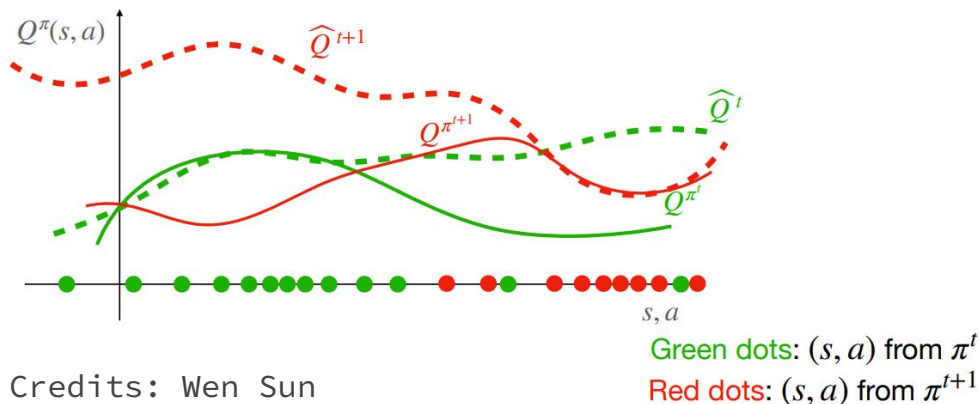


Credits: Wen Sun

Green dots:  $(s, a)$  from  $\pi^t$   
Red dots:  $(s, a)$  from  $\pi^{t+1}$



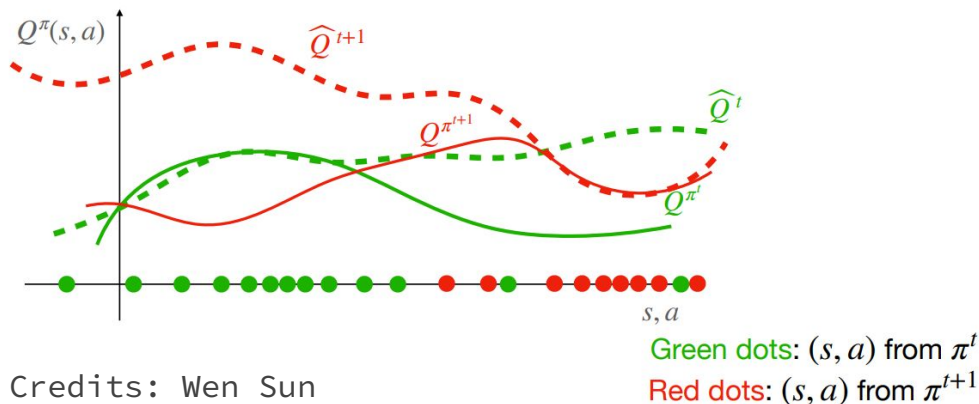
# Oscillation from Distribution Change



We cannot guarantee anymore monotonic improvement!



# Oscillation from Distribution Change



Our estimation is only good under  $d_{\mu_0}^\pi$  and to make sure we have monotonic improvement we need a strong coverage assumption

# Data Generation

— — —

2 steps:

1. Roll-in
2. Roll-out & compute supervision targets





# Data Generation

---

2 steps:

**1. Roll-in**

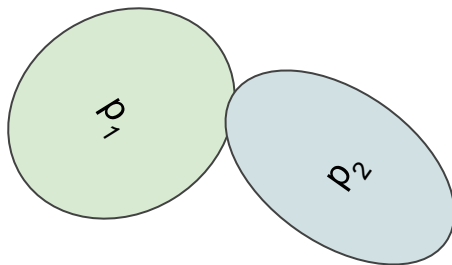
2. Roll-out & compute supervision targets

We want to sample our  $(s,a) \sim d^{\pi}_{s_0}(s,a) = (1-\gamma) \sum_{h=0}^{\infty} \gamma^h \mathbb{P}^{\pi}_h(s,a;s_0)$



# Sampling From Mixtures

— — —



$$p = (1-\alpha)p_1 + \alpha p_2$$

- Flip a coin with probability  $[\alpha, 1-\alpha]$
- Commit to a specific  $p_i$  based on that and sample from  $p_i$



# Data Generation

---

2 steps:

## 1. Roll-in

2. Roll-out & compute supervision targets

We want to sample our  $(s,a) \sim d^{\pi}_{s_0}(s,a) = (1-\gamma) \sum_{h=0}^{\infty} \gamma^h \mathbb{P}^{\pi_h}(s,a; s_0)$

- Sample  $h$  from  $\gamma^h(1-\gamma)$ , thus committing to a specific  $\mathbb{P}^{\pi_h}(s,a; s_0)$



# Data Generation

---

2 steps:

## 1. Roll-in

2. Roll-out & compute supervision targets

We want to sample our  $(s,a) \sim d^{\pi}_{s_0}(s,a) = (1-\gamma) \sum_{h=0}^{\infty} \gamma^h \mathbb{P}^{\pi}_h(s,a;s_0)$

- Sample  $h$  from  $\gamma^h(1-\gamma)$ , thus committing to a specific  $\mathbb{P}^{\pi}_h(s,a;s_0)$
- Follow  $\pi$  for  $h$  timesteps starting from  $s_0 \sim \mu_0$  and get  $s_h, a_h$



# Data Generation

— — —

2 steps:

1. Roll-in
- 2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?



# Data Generation

---

2 steps:

1. Roll-in

**2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

$$Q^\pi(s_t, a_t) = \mathbb{E}[\sum_{h=0}^{\infty} \gamma^h r_h \mid (s_0, a_0) = (s_t, a_t), a_{h+1} = \pi(s_h), s_{h+1} \sim p(\cdot \mid s_h, a_h)]$$

$$Q^\pi(s_t, a) = r_t + \gamma \mathbb{E}_{s' \sim p(\cdot \mid s, a)} [V^\pi(s')]$$



# Data Generation

---

2 steps:

1. Roll-in
- 2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

$$Q^\pi(s_t, a_t) = \mathbb{E}[\sum_{h=0}^{\infty} \gamma^h r_h \mid (s_0, a_0) = (s_t, a_t), a_{h+1} = \pi(s_h), s_{h+1} \sim p(\cdot \mid s_h, a_h)]$$

How do we get an unbiased  
estimate of this?



# Data Generation

---

2 steps:

1. Roll-in
- 2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

$$Q^\pi(s_t, a_t) = \mathbb{E}[\sum_{h=0}^{\infty} \gamma^h r_h | (s_0, a_0) = (s_t, a_t), a_{h+1} = \pi(s_h), s_{h+1} \sim p(\cdot | s_h, a_h)]$$

Sample many times and average!





# Data Generation

---

2 steps:

1. Roll-in

**2. Roll-out & compute supervision targets**


Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

$$Q^\pi(s_t, a_t) = \mathbb{E}[\sum_{h=0}^{\infty} \gamma^h r_h | (s_0, a_0) = (s_t, a_t), a_{h+1} = \pi(s_h), s_{h+1} \sim p(\cdot | s_h, a_h)]$$

Sample many times and average!



SAPIENZA  
UNIVERSITÀ DI ROMA



# Data Generation

---

2 steps:

1. Roll-in
- 2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

$$Q^\pi(s_t, a_t) = \mathbb{E}[\sum_{h=0}^{\infty} \gamma^h r_h \mid (s_0, a_0) = (s_t, a_t), a_{h+1} = \pi(s_h), s_{h+1} \sim p(\cdot \mid s_h, a_h)]$$

Easier said than done :(

**Infinite horizon!**



SAPIENZA  
UNIVERSITÀ DI ROMA

# Data Generation

---

2 steps:

1. Roll-in
- 2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

$$Q^\pi(s_t, a_t) = \mathbb{E}[\sum_{h=0}^{\infty} \gamma^h r_h | (s_0, a_0) = (s_t, a_t), a_{h+1} = \pi(s_h), s_{h+1} \sim p(\cdot | s_h, a_h)]$$

Use  $\gamma$  as a sampling factor  
again to choose an horizon



# Data Generation

---

2 steps:

1. Roll-in

**2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

- Start at  $s,a$
- Repeat:
  - Get  $r(s,a)$
  - With probability  $1-\gamma$  terminate and return  $y = \sum \gamma^h r_h$
  - Execute action and get in  $s'$



# Data Generation

---

2 steps:

1. Roll-in

**2. Roll-out & compute supervision targets**

Given  $s$ ,  $a$ , how do we estimate  $Q^\pi(s,a)$ ?

- Start at  $s,a$
- Repeat:
  - Get  $r(s,a)$
  - With probability  $1-\gamma$  terminate and return  $y = \sum \gamma^h r_h$
  - Execute action and get in  $s'$

$$D = \{s_i, a_i, y_i\}_{i=1}^N$$

