

Model-Based RL

Roberto Capobianco




SAPIENZA
UNIVERSITÀ DI ROMA

Recap

Deep Neural Networks

— — —

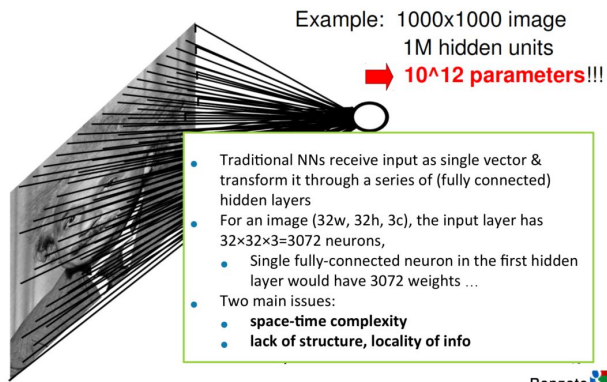


- Composition of multiple layers (functions)
- To fit the parameters, require a loss function (a measure of error)
- Use chain rule to propagate the error
- Combine linear and non-linear transformations

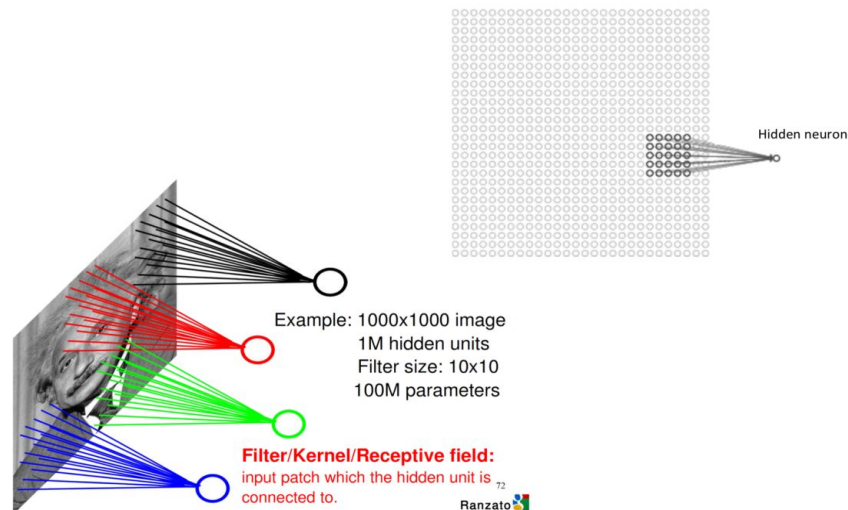

Why DNNs instead of feature engineering + other function approximators?

- Use distributed representations instead of local representations (like Kernels in SVMs)
- Universal function approximators
- Can potentially need exponentially less nodes/parameters (compared to a shallow net) to represent the same function
- Can learn the parameters using stochastic gradient descent

Convolutional Neural Networks



Ranzato



Ranzato




SAPIENZA
UNIVERSITÀ DI ROMA

Replay Buffer

— — —


- Used in most recent RL algorithms
 - Stores previous transitions experienced by an agent
 - Transitions are used later in time for training
-
- **Capacity:** total number of transitions stored in the buffer
 - **Age of transitions:** number of gradient steps taken by the learner since the transition was generated
 - Efficient use of previous experience
 - Same data used multiple times
 - Collecting real data costs
 - Better convergence with function approximators
 - Makes data more i.i.d.
 - More similar to a supervised learning task




DQN

s_1, a_1, r_2, s_2
s_2, a_2, r_3, s_3
s_3, a_3, r_4, s_4
...
$s_t, a_t, r_{t+1}, s_{t+1}$

→ s, a, r, s'



- End-to-end learning of values $Q(s, a)$ from pixels s
- Input state s is stack of raw pixels from last 4 frames
- Output is $Q(s, a)$ for 18 joystick/button positions
- Reward is change in score for that step



DQN Pseudo-code

```
1: Input  $\mathcal{C}$ ,  $\alpha$ ,  $D = \{\}$ , Initialize  $\mathbf{w}$ ,  $\mathbf{w}^- = \mathbf{w}$ ,  $t = 0$ 
2: Get initial state  $s_0$ 
3: loop
4:   Sample action  $a_t$  given  $\epsilon$ -greedy policy for current  $\hat{Q}(s_t, a; \mathbf{w})$ 
5:   Observe reward  $r_t$  and next state  $s_{t+1}$ 
6:   Store transition  $(s_t, a_t, r_t, s_{t+1})$  in replay buffer  $D$ 
7:   Sample random minibatch of tuples  $(s_i, a_i, r_i, s_{i+1})$  from  $D$ 
8:   for  $j$  in minibatch do
9:     if episode terminated at step  $i + 1$  then
10:       $y_i = r_i$ 
11:     else
12:       $y_i = r_i + \gamma \max_{a'} \hat{Q}(s_{i+1}, a'; \mathbf{w}^-)$ 
13:     end if
14:     Do gradient descent step on  $(y_i - \hat{Q}(s_i, a_i; \mathbf{w}))^2$  for parameters  $\mathbf{w}$ :  $\Delta \mathbf{w} = \alpha(y_i - \hat{Q}(s_i, a_i; \mathbf{w})) \nabla_{\mathbf{w}} \hat{Q}(s_i, a_i; \mathbf{w})$ 
15:   end for
16:    $t = t + 1$ 
17:   if  $\text{mod}(t, \mathcal{C}) == 0$  then
18:      $\mathbf{w}^- \leftarrow \mathbf{w}$ 
19:   end if
20: end loop
```

Credits: Emma Brunskill



DQN Result Analysis

Game	Linear	Deep Network	DQN w/ fixed Q	DQN w/ replay	DQN w/replay and fixed Q
Breakout	3	3	10	241	317
Enduro	62	29	141	831	1006
River Raid	2345	1453	2868	4102	7447
Seaquest	656	275	1003	823	2894
Space Invaders	301	302	373	826	1089

Credits: Emma Brunskill



SAPIENZA
UNIVERSITÀ DI ROMA

Maximization Bias

— — —

Greedy and epsilon greedy require a maximization step

This can lead to a positive bias:


- Consider s where many actions have $Q(s,a)$ that are all zero
- Estimated values are uncertain, some above and some below zero
- Maximum of the true values is zero, but max over estimates is positive

DDQN

- Extend Double Q-Learning to DQN
- Use current Q-network for action selection
- Use older (the target) Q-network (w^-) for action evaluation

$$\Delta \mathbf{w} = \alpha(r + \gamma \underbrace{\hat{Q}(\arg \max_{a'} \hat{Q}(s', a'; \mathbf{w}); \mathbf{w}^-)}_{\text{Action selection: } \mathbf{w}}) - \hat{Q}(s, a; \mathbf{w}))$$

Action evaluation: \mathbf{w}^-



Prioritized Experience Replay (PER)

— — —

Schaul, Tom, et al. "Prioritized experience replay." *arXiv preprint arXiv:1511.05952* (2015). - Code:

https://github.com/google/dopamine/tree/master/dopamine/replay_memory

- **Idea:** *“more frequently replay transitions with high expected learning progress, as measured by the magnitude of their temporal-difference (TD) error”*
- Prioritization can lead to:
 - Loss of diversity:
 - Use stochastic prioritization
 - Bias
 - Use importance sampling



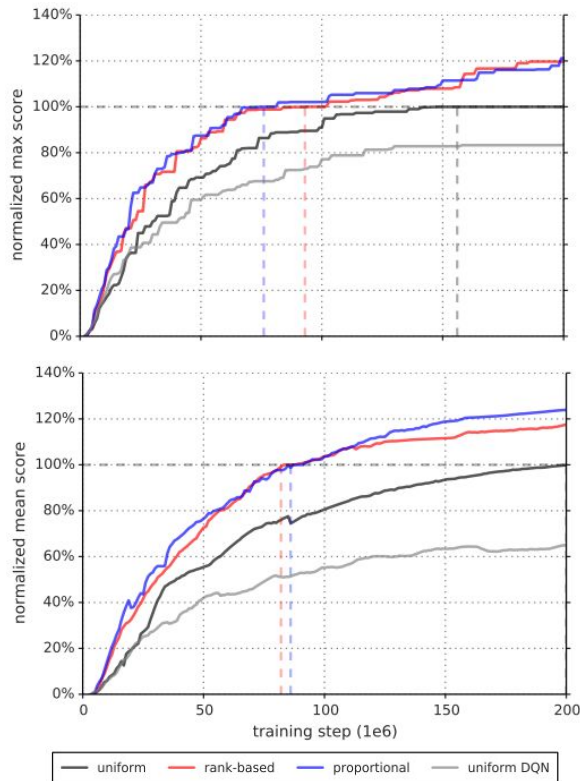
DDQN with PER

Algorithm 1 Double DQN with proportional prioritization

```

1: Input: minibatch  $k$ , step-size  $\eta$ , replay period  $K$  and size  $N$ , exponents  $\alpha$  and  $\beta$ , budget  $T$ .
2: Initialize replay memory  $\mathcal{H} = \emptyset$ ,  $\Delta = 0$ ,  $p_1 = 1$ 
3: Observe  $S_0$  and choose  $A_0 \sim \pi_\theta(S_0)$ 
4: for  $t = 1$  to  $T$  do
5:   Observe  $S_t, R_t, \gamma_t$ 
6:   Store transition  $(S_{t-1}, A_{t-1}, R_t, \gamma_t, S_t)$  in  $\mathcal{H}$  with maximal priority  $p_t = \max_{i < t} p_i$ 
7:   if  $t \equiv 0 \pmod K$  then
8:     for  $j = 1$  to  $k$  do
9:       Sample transition  $j \sim P(j) = p_j^\alpha / \sum_i p_i^\alpha$ 
10:      Compute importance-sampling weight  $w_j = (N \cdot P(j))^{-\beta} / \max_i w_i$ 
11:      Compute TD-error  $\delta_j = R_j + \gamma_j Q_{\text{target}}(S_j, \arg \max_a Q(S_j, a)) - Q(S_{j-1}, A_{j-1})$ 
12:      Update transition priority  $p_j \leftarrow |\delta_j|$ 
13:      Accumulate weight-change  $\Delta \leftarrow \Delta + w_j \cdot \delta_j \cdot \nabla_\theta Q(S_{j-1}, A_{j-1})$ 
14:    end for
15:    Update weights  $\theta \leftarrow \theta + \eta \cdot \Delta$ , reset  $\Delta = 0$ 
16:    From time to time copy weights into target network  $\theta_{\text{target}} \leftarrow \theta$ 
17:  end if
18:  Choose action  $A_t \sim \pi_\theta(S_t)$ 
19: end for

```



Importance Sampling for MC

The same idea is applied to RL for off-policy learning

Consider the MC setting: we want to use the returns from policy μ to evaluate π

Compute $G_t^{\pi/\mu}$ by multiplying **importance sampling corrections**

$$G_t^{\pi/\mu} = \frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} \frac{\pi(A_{t+1}|S_{t+1})}{\mu(A_{t+1}|S_{t+1})} \cdots \frac{\pi(A_T|S_T)}{\mu(A_T|S_T)} G_t$$

if μ is zero and π non-zero this cannot be used



Importance Sampling for TD

The same idea is applied to RL for off-policy learning

Consider the TD setting: we want to weight the TD target

$$\frac{\pi(A_t|S_t)}{\mu(A_t|S_t)} (R_{t+1} + \gamma V(S_{t+1}))$$

Q-Learning does not need it, why?

We directly use the action from the target policy




End Recap



SAPIENZA
UNIVERSITÀ DI ROMA

Model-Based RL: Motivation

— — —



Model-Based RL: Motivation

— — —



We cannot write out the exact analytical dynamics, but we can learn it from data $\{s, a, s'\}$




Model-Based RL: Motivation

— — —

We cannot write out the exact analytical
dynamics, but we can learn it from data
 $\{s, a, s'\}$

**And then find a policy by planning
on such dynamic model**




Model-Based RL

- **Model-free RL:** rely on learning alone
- **Model-based RL:** rely both on learning and planning on a *model*

Model: anything that an agent can use to predict how environment will respond to actions

- *Distribution models:* describe all possibilities and their probabilities
- *Sample models:* produce just one possibility sampled according to probabilities



Basic Algorithm

— — —

The simplest algorithm is the following:

1. Generate data


(e.g., execute a starting policy)

2. Fit a model using data

(e.g., using least-squares, or maximum likelihood)

3. Plan on the learned model

(e.g., using VI, PI, or LQR)



Basic Algorithm

— — —

The simplest algorithm is the following:

1. Generate data

(e.g., execute a starting policy)


2. Fit a model using data

(e.g., using least-squares, or maximum likelihood)

3. Plan on the learned model

(e.g., using VI, PI, or LQR)

Often iterate this process several times



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?



Simulation Lemma

approximation!

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?



Simulation Lemma

a model or simulator!

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

Example: if I have a driving policy and a self-driving car simulator: if I test such policy in the simulator and in the real world, what's the difference in terms of policy performance?



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in
the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in
the true dynamics



Simulation Lemma

— — —

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in
the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in
the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0)$$

?



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s,a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s,a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s,a)} \widehat{V}^{\pi}(s') \right]$$

distribution of s,a under the policy and the true dynamics



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

as usual, from the sum of the gammas



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

difference between the two distribution (model and real dynamics)



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

difference between two
distributions P and Q can be
computed as

$$|\mathbb{E}_{x \sim P}[f(x)] - \mathbb{E}_{x \sim Q}[f(x)]|$$

difference between the two distribution (model and real
dynamics)



SAPIENZA
UNIVERSITÀ DI ROMA

Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$



value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$


value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$


value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$



the action distribution is the
same, the only difference, at one
step, is in the next state



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$


value of policy in the true
dynamics

— — —

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

Computing this is very difficult,
but we can do it one step at a time



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator


$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

Let's step in the real dynamics for one step, and then go back to the simulator



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator


$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

We can do recursion and follow the
same reasoning again



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —
Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$

$$V^{\pi}(s_0) = r(s_0, \pi(s_0)) + \gamma \mathbb{E}_{s_1 \sim P(s_0, \pi(s_0))} [V^{\pi}(s_1)]$$

and similar for the \widehat{V} , so $r(s_0, \pi(s_0))$ cancels out



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$

$$= \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) + \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$

$$= \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) + \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$

$$= \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) \right] + \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0), s_1 \sim P(s_0, a_0)} \left[\widehat{V}^{\pi}(s_1) - V^{\pi}(s_1) \right]$$

Just play with the formula and
re-order stuff



SAPIENZA
UNIVERSITÀ DI ROMA

Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\boxed{\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0)} = \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$

$$= \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) + \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} V^{\pi}(s_1) \right]$$

$$= \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0)} \left[\mathbb{E}_{s_1 \sim \widehat{P}(s_0, a_0)} \widehat{V}^{\pi}(s_1) - \mathbb{E}_{s_1 \sim P(s_0, a_0)} \widehat{V}^{\pi}(s_1) \right] + \gamma \mathbb{E}_{a_0 \sim \pi(\cdot | s_0), s_1 \sim P(s_0, a_0)} \boxed{\widehat{V}^{\pi}(s_1) - V^{\pi}(s_1)}$$



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —
Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

...and we get to this just
by doing an exponential
averaging over all the
timesteps



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —
Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

We also know that

$$|\mathbb{E}_{x \sim P}[f(x)] - \mathbb{E}_{x \sim Q}[f(x)]| \leq \sup_x f(x) \|P - Q\|_1$$

where we use the l1-norm and the *total variation distance* between the two distributions



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\leq \frac{1}{(1 - \gamma)^2} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left\| \widehat{P}(\cdot \mid s, a) - P(\cdot \mid s, a) \right\|_1$$

Since we assume reward is in $[0, 1]$, the max value of V is $1/(1-\gamma)$



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —
Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?



$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\leq \frac{1}{(1 - \gamma)^2} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left\| \widehat{P}(\cdot \mid s, a) - P(\cdot \mid s, a) \right\|_1$$

We can bound the policy performance difference by the total model disagreement measured on the real trajectory



Full Model-Based RL Loop



Model Fitting

How can we fit a model?

For example, very simply, collect N data-points and estimate it as follows (note that we're using the indicator function **1**)

$$\widehat{P}(s' | s, a) = \frac{\sum_{i=1}^N \mathbf{1}\{s'_i = s'\}}{N}$$



Model Fitting

How can we fit a model?

For example, very simply, collect N data-points and estimate it as follows (note that we're using the indicator function **1**)

$$\widehat{P}(s' | s, a) = \frac{\sum_{i=1}^N \mathbf{1}\{s'_i = s'\}}{N}$$

At infinity this should converge to the true P



Planning

— — —

How can we plan using a model?

Use value iteration, policy iteration, LQR if we're in continuous space, or other solutions like:

- Q-planning
- Monte-Carlo Tree Search



Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

- $S \leftarrow$ current (nonterminal) state
- $A \leftarrow \varepsilon$ -greedy(S, Q)
- Take action A ; observe resultant reward, R , and state, S'
- $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
- $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)
- Loop repeat n times:


$S \leftarrow$ random previously observed state

$A \leftarrow$ random action previously taken in S



$R, S' \leftarrow Model(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)



Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

- $S \leftarrow$ current (nonterminal) state
- $A \leftarrow \varepsilon$ -greedy(S, Q)
- Take action A ; observe resultant reward, R , and state, S'
- $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
- $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)
- Loop repeat n times:

$S \leftarrow$ random previously observed state


$A \leftarrow$ random action previously taken in S

$R, S' \leftarrow Model(S, A)$



$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

Model Fitting

WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)



SAPIENZA
UNIVERSITÀ DI ROMA

Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

- $S \leftarrow$ current (nonterminal) state
- $A \leftarrow \varepsilon$ -greedy(S, Q)
- Take action A ; observe resultant reward, R , and state, S'
- $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
- $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)
- Loop repeat n times:

$S \leftarrow$ random previously observed state


$A \leftarrow$ random action previously taken in S

$R, S' \leftarrow Model(S, A)$

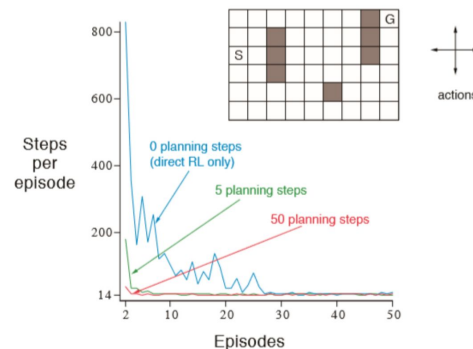

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

Q-planning

WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)



SAPIENZA
UNIVERSITÀ DI ROMA

More on Planning

— — —

Background planning (e.g., Dyna)

- Not focused on current state
- Gradually improve policy on the basis of simulated experience from model
- Planning plays a part well before an action is selected

Decision-time planning

- Begin planning after encountering each new state
- Evaluates action choices leading to different predicted states
- Use simulated experience to select an action for the current state
- Values and policy are updated specifically for current state

Can be blended together




Rollout Planning Algorithm

Decision-time planning

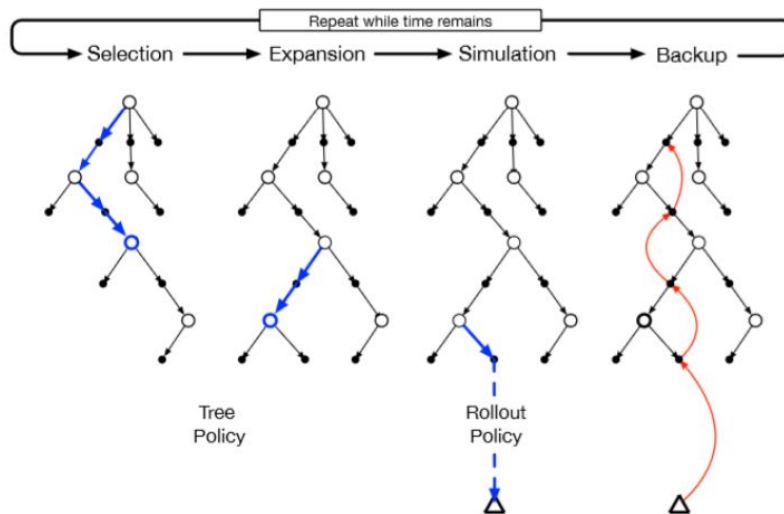
- Uses MC control applied to simulated trajectories starting at current state
- Estimate action values by averaging returns of many simulated trajectories: try each possible action for one step and then follow rollout policy
- When estimate accurate, highest value action is executed

Does not estimate (unlike MC) full value-function, but only value of actions for current state and given policy



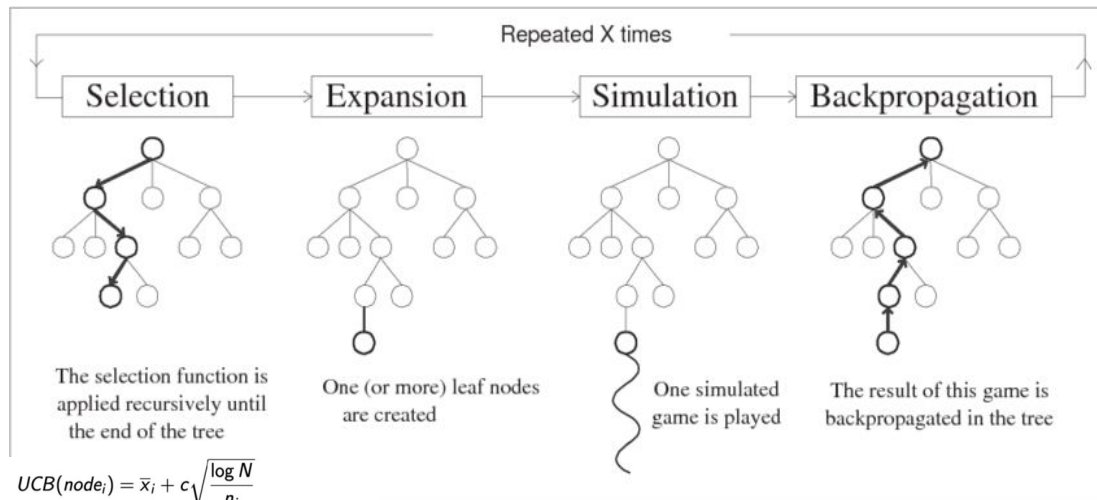
Monte-Carlo Tree Search

Decision-time planning, like a rollout algorithm BUT accumulating value estimates



Monte-Carlo Tree Search

Decision-time planning, like a rollout algorithm BUT accumulating value estimates



\bar{x}_i : mean node value; n_i : #visits of node i ; N #visits parent;



AlphaGo

— — —

Silver, David, et al. "Mastering the game of Go with deep neural networks and tree search." *nature* 529.7587 (2016): 484-489.



SAPIENZA
UNIVERSITÀ DI ROMA

AlphaGo - Networks

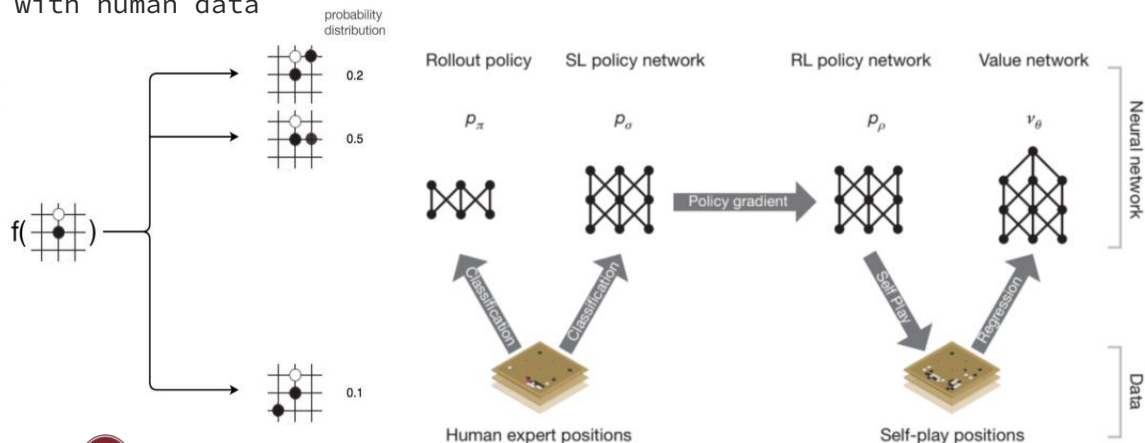
- Faster than SL Policy
- Less accurate than SL Policy
- Used during simulation

- 19×19×48 input feature to represent the board
- Trained on millions of board positions with human data


- RL policy initialized at SL policy
- RL training occurs using outcomes z (+1 win, -1 lose)

Extended Data Table 2 | Input features for neural networks

Feature	# of planes	Description
Stone colour	3	Player stone / opponent stone / empty
Ones	1	A constant plane filled with 1
Turns since	8	How many turns since a move was played
Liberties	8	Number of liberties (empty adjacent points)
Capture size	8	How many opponent stones would be captured
Self-atari size	8	How many of own stones would be captured
Liberties after move	8	Number of liberties after this move is played
Ladder capture	1	Whether a move at this point is a successful ladder capture
Ladder escape	1	Whether a move at this point is a successful ladder escape
Sensibleness	1	Whether a move is legal and does not fill its own eyes
Zeros	1	A constant plane filled with 0
Player color	1	Whether current player is black



AlphaGo - MCTS


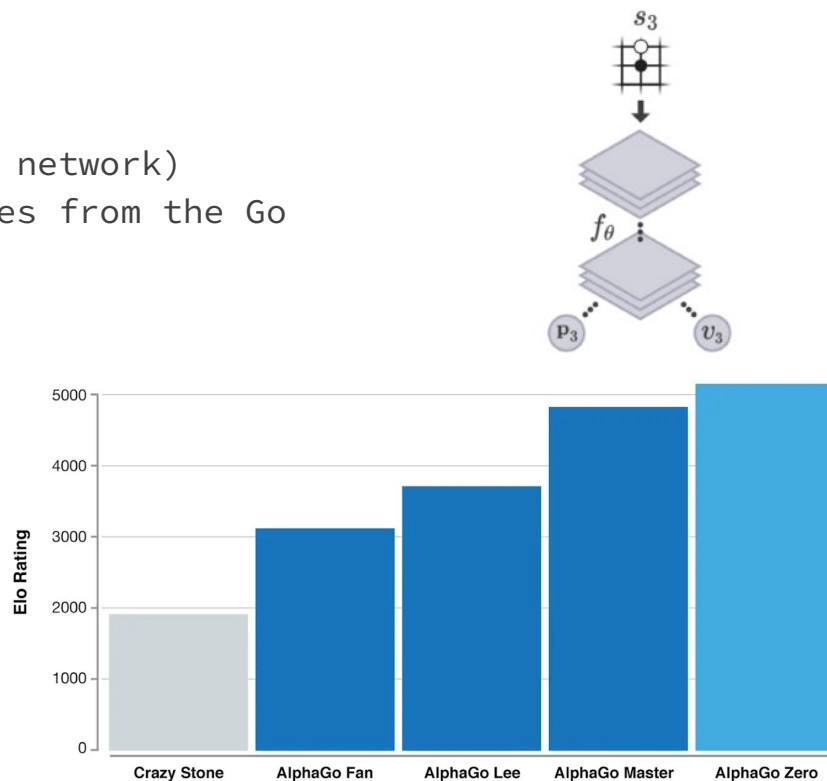



- Selection makes use of UCT
- Rollouts completed using rollout policy
- Backup done by mixing value network prediction and outcome of a rollout
- AlphaGo plays against previous versions of itself (self-play)




AlphaGo Zero VS AlphaGo

- Does not use human knowledge (SL policy network)
- Trained using only black and white stones from the Go board as input
- Uses one neural network rather than two
- Does not use rollouts




AlphaGo Zero VS AlphaGo - MCTS

AlphaGo Zero




AlphaGo




$$(p, v) = f_\theta \left(\begin{array}{c} \text{Board State} \\ \text{Action} \end{array} \right)$$

$(0.5, 0.45, \dots, 0.9)$




AlphaGo Zero VS AlphaGo - Training

a. Self-Play



b. Neural Network Training





MuZero - MCTS

— — —


Models:

- policy
- value function
- reward
- new hidden state



MuZero - Training

Learned model unrolled with the collected experience



Learning Dynamics from Pixels

— — —

Hafner, Danijar, et al. "Learning latent dynamics for planning from pixels." *International Conference on Machine Learning*. PMLR, 2019.

- Not always the state is available (POMDP): learn a compact representation
 - A recurrent model is needed
- Plan in the learned (latent!) dynamics space

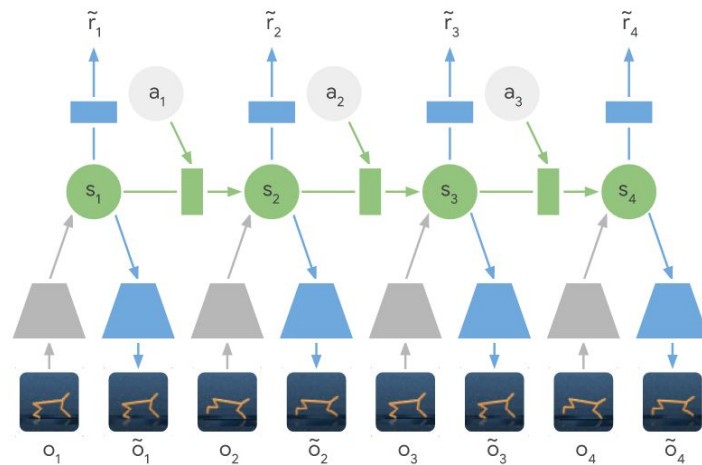
Transition function:	$s_t \sim p(s_t \mid s_{t-1}, a_{t-1})$	
Observation function:	$o_t \sim p(o_t \mid s_t)$	(1)
Reward function:	$r_t \sim p(r_t \mid s_t)$	
Policy:	$a_t \sim p(a_t \mid o_{\leq t}, a_{< t}),$	



PlaNet

— — —

- Learns a
 - transition model
 - observation model
 - reward model
- Policy obtained using planning in MPC fashion



PlaNet - Data Collection & Training

Algorithm 1: Deep Planning Network (PlaNet)

Input:

R Action repeat	$p(s_t s_{t-1}, a_{t-1})$	Transition model
S Seed episodes	$p(o_t s_t)$	Observation model
C Collect interval	$p(r_t s_t)$	Reward model
B Batch size	$q(s_t o_{\leq t}, a_{< t})$	Encoder
L Chunk length	$p(\epsilon)$	Exploration noise
α Learning rate		


```

1 Initialize dataset  $\mathcal{D}$  with  $S$  random seed episodes.
2 Initialize model parameters  $\theta$  randomly.
3 while not converged do
    // Model fitting
4   for update step  $s = 1..C$  do
5     Draw sequence chunks  $\{(o_t, a_t, r_t)_{t=k}^{L+k} \}_{i=1}^B \sim \mathcal{D}$ 
      uniformly at random from the dataset.
6     Compute loss  $\mathcal{L}(\theta)$  from Equation 8.
7     Update model parameters  $\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}(\theta)$ .

    // Data collection
8    $o_1 \leftarrow \text{env.reset}()$ 
9   for time step  $t = 1..\lceil \frac{T}{R} \rceil$  do
10    Infer belief over current state  $q(s_t | o_{\leq t}, a_{< t})$  from
      the history.
11     $a_t \leftarrow \text{planner}(q(s_t | o_{\leq t}, a_{< t}), p)$ , see
      Algorithm 2 in the appendix for details.
12    Add exploration noise  $\epsilon \sim p(\epsilon)$  to the action.
13    for action repeat  $k = 1..R$  do
14       $r_t^k, o_{t+1}^k \leftarrow \text{env.step}(a_t)$ 
15       $r_t, o_{t+1} \leftarrow \sum_{k=1}^R r_t^k, o_{t+1}^k$ 
16     $\mathcal{D} \leftarrow \mathcal{D} \cup \{(o_t, a_t, r_t)_{t=1}^T\}$ 

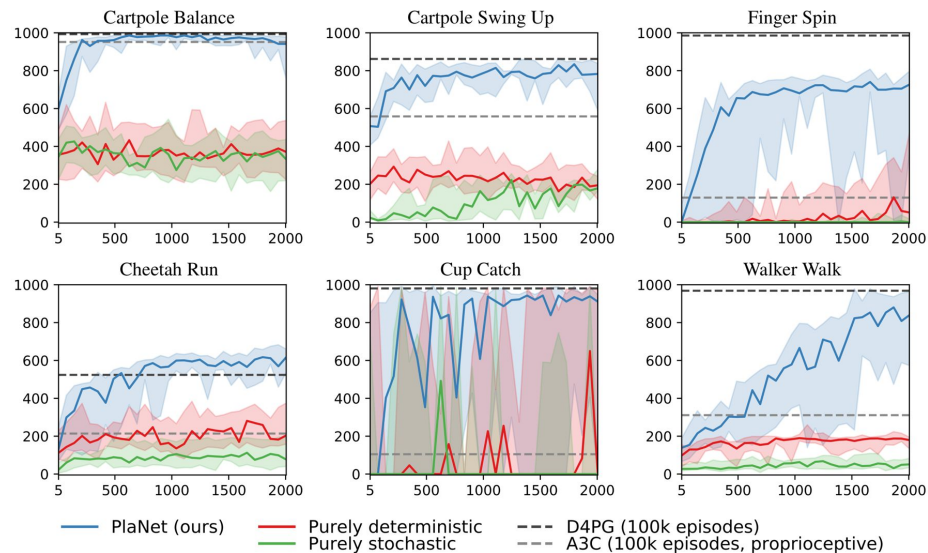
```

- Planning using Cross-Entropy Method
 - Population-based optimization algorithm that infers a distribution over action sequences that maximize the objective
- Encode past images
- Execute only first planned action



PlaNet Results

— — —



Sim-to-Real

— — —

Tan, Jie, et al. "Sim-to-real: Learning agile locomotion for quadruped robots." *arXiv preprint arXiv:1804.10332* (2018).



Narrowing the Reality Gap

— — —


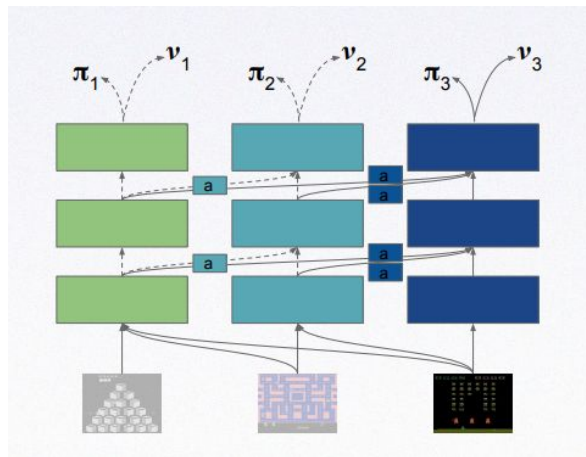
- Fine-tuning
- Progressive nets
- Improved modeling
 - System identification, better models, etc.
- Randomization
 - Random perturbations, Randomization in dynamics parameters



Progressive Networks


Rusu, Andrei A., et al. "Sim-to-real robot learning from pixels with progressive nets." arXiv preprint arXiv:1610.04286 (2016).

- Avoid catastrophic forgetting
- Allow transfer and multi-task
- Can be used to transfer from simulation to real robot



Progressive Nets VS Fine-Tuning

— — —



Dynamics Randomization

— — —

Comparisons



our method



no randomization
during training

