

Policy Search

Roberto Capobianco



SAPIENZA
UNIVERSITÀ DI ROMA


Recap

Model-Based RL: Motivation

— — —

We cannot write out the exact analytical
dynamics, but we can learn it from data
 $\{s, a, s'\}$

**And then find a policy by planning
on such dynamic model**



Basic Algorithm

— — —

The simplest algorithm is the following:

1. Generate data

(e.g., execute a starting policy)


2. Fit a model using data

(e.g., using least-squares, or maximum likelihood)

3. Plan on the learned model

(e.g., using VI, PI, or LQR)

Often iterate this process several times



Simulation Lemma

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true dynamics

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics


Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?

$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

Computing this is very difficult,
but we can do it one step at a
time.

At a single step the action
distribution is the same, the only
difference, is in the next state:
Let's step in the real dynamics for
one step, and then go back to the
simulator.

We can do recursion and follow the
same reasoning again.



Simulation Lemma

$$\widehat{V}^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, \widehat{P} \right]$$

value of policy in the
simulator

$$V^{\pi}(s_0) = \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid \pi, P \right]$$

value of policy in the true
dynamics

— — —

Question: If I have an estimator of the true dynamics, what's the performance of a policy under this estimator, wrt the real world?



$$\widehat{V}^{\pi}(s_0) - V^{\pi}(s_0) = \frac{\gamma}{1 - \gamma} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left[\mathbb{E}_{s' \sim \widehat{P}(s, a)} \widehat{V}^{\pi}(s') - \mathbb{E}_{s' \sim P(s, a)} \widehat{V}^{\pi}(s') \right]$$

$$\leq \frac{1}{(1 - \gamma)^2} \mathbb{E}_{s, a \sim d_{s_0}^{\pi}} \left\| \widehat{P}(\cdot \mid s, a) - P(\cdot \mid s, a) \right\|_1$$

We can bound the policy performance difference by the total model disagreement measured on the real trajectory



Full Model-Based RL Loop



Model Fitting

— — —

How can we fit a model?

For example, very simply, collect N data-points and estimate it as follows (note that we're using the indicator function $\mathbf{1}$)

$$\widehat{P}(s'|s, a) = \frac{\sum_{i=1}^N \mathbf{1}\{s'_i = s'\}}{N}$$

At infinity this should converge to the true P

How can we plan using a model?

Use value iteration, policy iteration, LQR if we're in continuous space, or other solutions like:

- Q-planning
- Monte-Carlo Tree Search



Dyna-Q

Initialize $Q(s, a)$ and $Model(s, a)$ for all $s \in \mathcal{S}$ and $a \in \mathcal{A}(s)$

Loop forever:

- $S \leftarrow$ current (nonterminal) state
- $A \leftarrow \varepsilon$ -greedy(S, Q)
- Take action A ; observe resultant reward, R , and state, S'
- $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
- $Model(S, A) \leftarrow R, S'$ (assuming deterministic environment)
- Loop repeat n times:


$S \leftarrow$ random previously observed state

$A \leftarrow$ random action previously taken in S



$R, S' \leftarrow Model(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

WITHOUT PLANNING ($n=0$)



WITH PLANNING ($n=50$)




Rollout Planning Algorithm

Decision-time planning

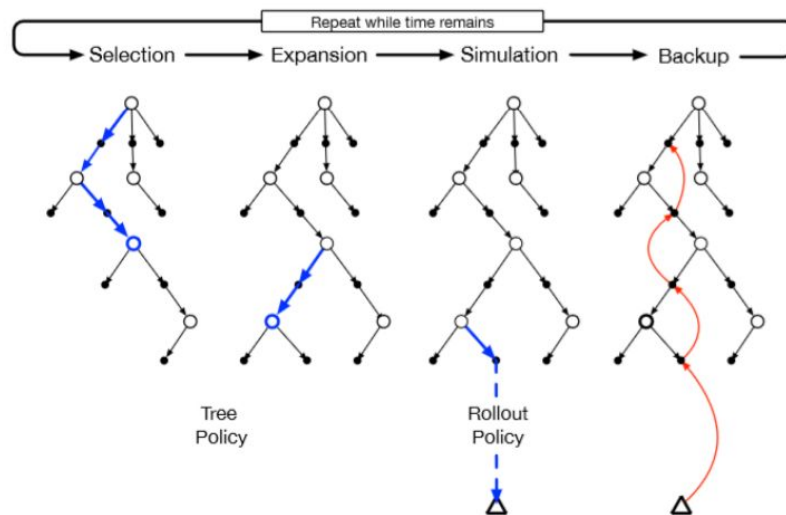
- Uses MC control applied to simulated trajectories starting at current state
- Estimate action values by averaging returns of many simulated trajectories: try each possible action for one step and then follow rollout policy
- When estimate accurate, highest value action is executed


Does not estimate (unlike MC) full value-function, but only value of actions for current state and given policy



Monte-Carlo Tree Search

Decision-time planning, like a rollout algorithm BUT accumulating value estimates





Learning Dynamics from Pixels

— — —

Hafner, Daniyar, et al. "Learning latent dynamics for planning from pixels." *International Conference on Machine Learning*. PMLR, 2019.

- Not always the state is available (POMDP): learn a compact representation
 - A recurrent model is needed
- Plan in the learned (latent!) dynamics space

| | | |
|-----------------------|---|-----|
| Transition function: | $s_t \sim p(s_t \mid s_{t-1}, a_{t-1})$ | |
| Observation function: | $o_t \sim p(o_t \mid s_t)$ | (1) |
| Reward function: | $r_t \sim p(r_t \mid s_t)$ | |
| Policy: | $a_t \sim p(a_t \mid o_{\leq t}, a_{< t}),$ | |



Narrowing the Reality Gap

— — —

- Fine-tuning
- Progressive nets
- Improved modeling
 - System identification, better models, etc.
- Randomization
 - Random perturbations, Randomization in dynamics parameters



End Recap



SAPIENZA
UNIVERSITÀ DI ROMA

Policy from Value

So far, we derived a policy out of a tabular or parameterized state-action value function



Policy from Value: Problems (1)

So far, we derived a policy out of a tabular or parameterized state-action value function

Computing this is very difficult if the action space is large



Policy from Value: Problems (2)

So far, we derived a policy out of a tabular or parameterized state-action value function

Reward: -1 per step

Actions: left, right

Action effects are as usual
in first and third states,
reversed in second state



All states appear identical
in their featurization

$x(s, \text{right}) = [1, 0]$ and


$x(s, \text{left}) = [0, 1]$



Policy from Value: Problems (2)

So far, we derived a policy out of a tabular or parameterized state-action value function


The optimal policy is not always deterministic and eps-greedy is not enough



Policy from Value: Problems (3)

So far, we derived a policy out of a tabular or parameterized state-action value function

Dramatic policy oscillations due to small value changes



Credits: Wen Sun



SAPIENZA
UNIVERSITÀ DI ROMA

Parameterized Policy

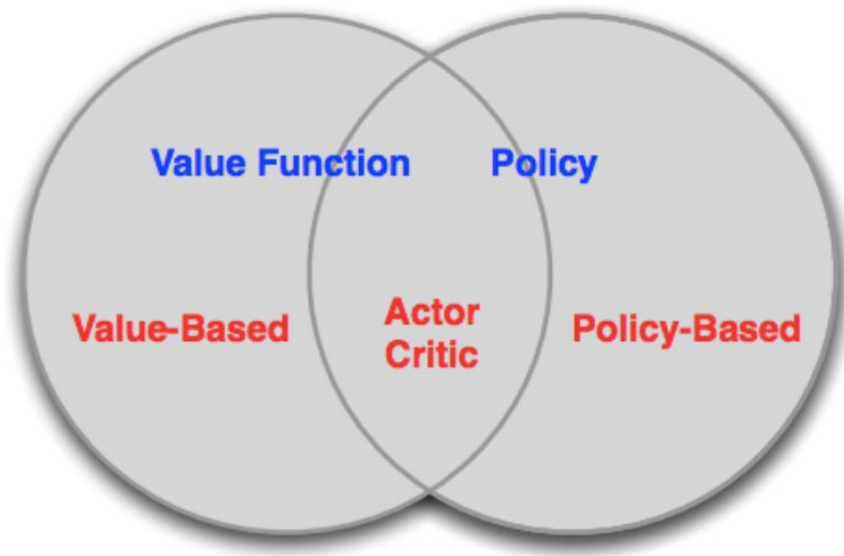
Can we directly represent and learn a parameterized policy?

$$a \sim \pi_{\theta}(s) \text{ or } p(\cdot | s, \theta)$$



Policy-Based RL

— — —



Credits: David Silver



SAPIENZA
UNIVERSITÀ DI ROMA

Finding a Parameterized Policy

How do we search for a parameterized policy $\pi_{\theta}(s)$?

We want to find the parameters θ such that π_{θ} is the best policy



Finding a Parameterized Policy

How do we search for a parameterized policy $\pi_{\theta}(s)$?

We want to find the parameters θ such that π_{θ} is the best policy

How do we measure the quality for a policy?



Quality of a Parameterized Policy

— — —

How do we measure the quality for a policy?

In other words, what is our objective function to maximize through the policy parameters?

Quality of a Parameterized Policy

How do we measure the quality for a policy?

In other words, what is our objective function to maximize through the policy parameters?

- For episodic tasks:
 - Value at start state (undiscounted)
- For continuing tasks:
 - Value at start state (discounted)
 - Average value

$$J(\pi) := \mathbb{E} \left[\sum_{h=0}^{H-1} r(s_h, a_h) \mid s_0 \sim \mu, s_{h+1} \sim P_{s_h, a_h}, a_h \sim \pi(\cdot \mid s_h) \right]$$

$$J(\pi) := \mathbb{E} \left[\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h) \mid s_0 \sim \mu, s_{h+1} \sim P_{s_h, a_h}, a_h \sim \pi(\cdot \mid s_h) \right]$$

see [Sutton&Barto 10.4: Deprecating the Discounted Setting]

- Average reward per timestep

see [Sutton&Barto 10.4: Deprecating the Discounted Setting]



Policy Optimization

Can use gradient free optimization

- Hill climbing Simplex / Nelder Mead Genetic algorithms
- Cross-Entropy method (CEM)
- Covariance Matrix Adaptation (CMA)



Policy Optimization

Can use gradient free optimization

- Hill climbing Simplex / Nelder Mead Genetic algorithms
- Cross-Entropy method (CEM)
- Covariance Matrix Adaptation (CMA)

Evolution strategies can rival the performance of standard RL techniques on modern RL benchmarks (e.g. Atari/MuJoCo):
<https://openai.com/blog/evolution-strategies/>



Policy Optimization

Can use gradient free optimization

- Hill climbing Simplex / Nelder Mead Genetic algorithms
- Cross-Entropy method (CEM)
- Covariance Matrix Adaptation (CMA)


Can work with any policy parameterizations, including non-differentiable ones, but can be sample inefficient and requires higher outcome variability



Policy Gradient

— — —

An efficient and widely used technique to find a parameterized policy is the policy gradient



Policy Gradient

An efficient and widely used technique to find a parameterized policy is the policy gradient

Guaranteed to converge to local maximum or global maximum, but it often converges only to a local maximum



Policy Gradient

An efficient and widely used technique to find a parameterized policy is the policy gradient

Only for convex function,
this guarantees global
optimality



Policy Gradient

$$\pi_{\theta}(a | s) = \pi(a | s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta}) |_{\theta=\theta_t}$$



Policy Gradient

$$\pi_{\theta}(a | s) = \pi(a | s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \boxed{\nabla_{\theta} J(\pi_{\theta}) |_{\theta=\theta_t}}$$

HOW?



Finite Differences

— — —

$$\pi_{\theta}(a | s) = \pi(a | s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

$$J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$



Finite Differences

— — —

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

$$J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$



Finite Differences

— — —

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

$$J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$V^{\pi}(s_{\theta}) !$$



Finite Differences

— — —

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

$$J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$V^{\pi}(s_{\theta}) !$$



Finite Differences

— — —

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

IDEA: perturb θ by small amount in k -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon u_k) - V(s_0, \theta)}{\epsilon}$$



Finite Differences

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

IDEA: perturb θ by small amount in k -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon \mathbf{u}_k) - V(s_0, \theta)}{\epsilon} \quad \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \leftarrow \begin{matrix} k\text{-th} \\ \text{entry} \end{matrix}$$



Finite Differences

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

IDEA: perturb θ by small amount in k -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon \mathbf{u}_k) - V(s_0, \theta)}{\epsilon} \quad \begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \leftarrow \begin{matrix} k\text{-th} \\ \text{entry} \end{matrix}$$

Uses n evaluations to compute policy gradient in n dimensions



Finite Differences

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

IDEA: perturb θ by small amount in k -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon \mathbf{u}_k) - V(s_0, \theta)}{\epsilon}$$

$$\begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \leftarrow \begin{matrix} k\text{-th} \\ \text{entry} \end{matrix}$$

Uses n evaluations to compute policy gradient in n dimensions

Works for arbitrary policies, even if policy is not differentiable



Finite Differences

$$\pi_{\theta}(a|s) = \pi(a|s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta})|_{\theta=\theta_t}$$

Simple approach: compute $\nabla_{\theta} J(\pi_{\theta})$ using finite differences

IDEA: perturb θ by small amount in k -th dimension

$$\frac{\partial V(s_0, \theta)}{\partial \theta_k} \approx \frac{V(s_0, \theta + \epsilon \mathbf{u}_k) - V(s_0, \theta)}{\epsilon}$$

$$\begin{pmatrix} 0 \\ 0 \\ \vdots \\ 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{pmatrix} \leftarrow \begin{matrix} k\text{-th} \\ \text{entry} \end{matrix}$$

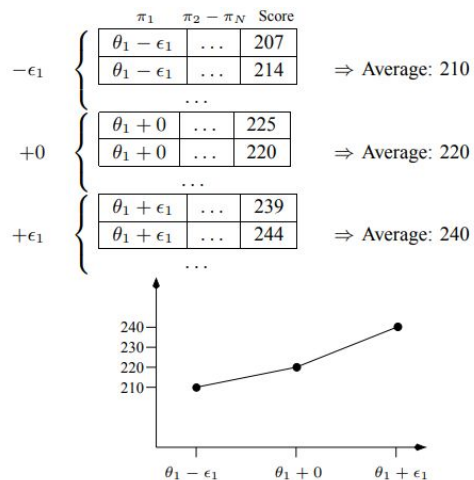
Uses n evaluations to compute policy gradient in n dimensions

Noise can dominate, but can be reduced by averaging over many samples or by reducing randomness if possible



Finite Differences: Example

Policy Gradient Reinforcement Learning for Fast Quadrupedal Locomotion. Nate Kohl and Peter Stone. Proceedings of the IEEE International Conference on Robotics and Automation, pp. 2619--2624, May 2004.



```

 $\pi \leftarrow \text{InitialPolicy}$ 
while !done do
   $\{R_1, R_2, \dots, R_t\} = t$  random perturbations of  $\pi$ 
  evaluate(  $\{R_1, R_2, \dots, R_t\}$  )
  for  $n = 1$  to  $N$  do
     $Avg_{+,n} \leftarrow$  average score for all  $R_i$  that have a positive
      perturbation in dimension  $n$ 
     $Avg_{+,n} \leftarrow$  average score for all  $R_i$  that have a zero
      perturbation in dimension  $n$ 
     $Avg_{-,n} \leftarrow$  average score for all  $R_i$  that have a
      negative perturbation in dimension  $n$ 
    if  $Avg_{+,n} > Avg_{+,n}$  and  $Avg_{+,n} > Avg_{-,n}$  then
       $A_n \leftarrow 0$ 
    else
       $A_n \leftarrow Avg_{+,n} - Avg_{-,n}$ 
    end if
  end for
   $A \leftarrow \frac{A}{|A|} * \eta$ 
   $\pi \leftarrow \pi + A$ 
end while
  
```



Policy Gradient

$$\pi_{\theta}(a | s) = \pi(a | s; \theta) \quad J(\pi_{\theta}) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{h=0}^{\infty} \gamma^h r_h \right]$$

$$\theta_{t+1} = \theta_t + \eta \nabla_{\theta} J(\pi_{\theta}) |_{\theta=\theta_t}$$

Can we compute the gradient analytically?



Differentiable Policy Classes

— — —

Softmax Linear Policy

$$\pi_{\theta}(a | s) = \frac{\exp(\theta^{\top} \phi(s, a))}{\sum_{a'} \exp(\theta^{\top} \phi(s, a'))}$$



Differentiable Policy Classes

— — —

Softmax Linear Policy

$$\pi_{\theta}(a | s) = \frac{\exp(\theta^{\top} \phi(s, a))}{\sum_{a'} \exp(\theta^{\top} \phi(s, a'))}$$

feature vector



Differentiable Policy Classes

— — —

Softmax Linear Policy

$$\pi_{\theta}(a | s) = \frac{\exp(\theta^{\top} \phi(s, a))}{\sum_{a'} \exp(\theta^{\top} \phi(s, a'))}$$



parameters



Differentiable Policy Classes

Softmax Linear Policy

$$\pi_{\theta}(a | s) = \frac{\exp(\theta^{\top} \phi(s, a))}{\sum_{a'} \exp(\theta^{\top} \phi(s, a'))}$$



Can be thought of as a classifier for discrete actions



Differentiable Policy Classes

— — —

Softmax Linear Policy

$$\pi_{\theta}(a | s) = \frac{\exp(\theta^{\top} \phi(s, a))}{\sum_{a'} \exp(\theta^{\top} \phi(s, a'))}$$

Softmax Policy

$$\pi_{\theta}(a | s) = \frac{\exp(f_{\theta}(s, a))}{\sum_{a'} \exp(f_{\theta}(s, a'))}$$

f can be a neural network



Differentiable Policy Classes

— — —

Softmax Linear Policy

$$\pi_{\theta}(a | s) = \frac{\exp(\theta^{\top} \phi(s, a))}{\sum_{a'} \exp(\theta^{\top} \phi(s, a'))}$$

Softmax Policy

$$\pi_{\theta}(a | s) = \frac{\exp(f_{\theta}(s, a))}{\sum_{a'} \exp(f_{\theta}(s, a'))}$$

Gaussian Policy

$$\pi_{\theta}(a | s) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(a - f_{\theta}(s))^2}{2\sigma^2}\right)$$

f can be a neural network



Likelihood Ratio Policy Gradient

Suppose we have a trajectory $\tau = \{s_0, a_0, s_1, a_1, \dots\}$

it's probability distribution, depending on θ is

$$\rho_{\theta}(\tau) = \mu(s_0)\pi_{\theta}(a_0 | s_0)P(s_1 | s_0, a_0)\pi_{\theta}(a_1 | s_1)\dots$$



Likelihood Ratio Policy Gradient

Suppose we have a trajectory $\tau = \{s_0, a_0, s_1, a_1, \dots\}$
it's probability distribution, depending on θ is

$$\rho_{\theta}(\tau) = \mu(s_0)\pi_{\theta}(a_0 | s_0)P(s_1 | s_0, a_0)\pi_{\theta}(a_1 | s_1)\dots$$

$$J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \rho_{\theta}(\tau)} \left[\underbrace{\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h)}_{R(\tau)} \right] = \sum_{\tau} P(\tau; \theta) R(\tau)$$



Likelihood Ratio Policy Gradient

Suppose we have a trajectory $\tau = \{s_0, a_0, s_1, a_1, \dots\}$

it's probability distribution, depending on θ is

$$\rho_{\theta}(\tau) = \mu(s_0)\pi_{\theta}(a_0 | s_0)P(s_1 | s_0, a_0)\pi_{\theta}(a_1 | s_1)\dots$$

$$J(\pi_{\theta}) = \mathbb{E}_{\tau \sim \rho_{\theta}(\tau)} \left[\underbrace{\sum_{h=0}^{\infty} \gamma^h r(s_h, a_h)}_{R(\tau)} \right] = \sum_{\tau} \boxed{P(\tau; \theta)} R(\tau)$$

probability of
each trajectory
from the
distribution



Likelihood Ratio Policy Gradient

We can then rewrite our objective as

$$\arg \max_{\theta} J(\pi_{\theta}) = \arg \max_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau)$$

Likelihood Ratio Policy Gradient

We can then rewrite our objective as

$$\arg \max_{\theta} J(\pi_{\theta}) = \arg \max_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau)$$



$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau)$$

Likelihood Ratio Policy Gradient

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} \nabla_{\theta} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} \frac{P(\tau; \theta)}{P(\tau; \theta)} \nabla_{\theta} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} P(\tau; \theta) \frac{\nabla_{\theta} P(\tau; \theta)}{P(\tau; \theta)} R(\tau) \\ &= \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau)\end{aligned}$$

Likelihood Ratio Policy Gradient

$$\begin{aligned}\nabla_{\theta} J(\theta) &= \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} \nabla_{\theta} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} \frac{P(\tau; \theta)}{P(\tau; \theta)} \nabla_{\theta} P(\tau; \theta) R(\tau) \\ &= \sum_{\tau} P(\tau; \theta) \frac{\nabla_{\theta} P(\tau; \theta)}{P(\tau; \theta)} R(\tau) \\ &= \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau)\end{aligned}$$

Can also be
derived/generalized
through an importance
sampling derivation

Likelihood Ratio Policy Gradient

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) = \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau)$$

Estimate this empirically as

$$\approx \hat{g} = (1/m) \sum_{i=1}^m R(\tau^{(i)}) \nabla_{\theta} \log P(\tau^{(i)}; \theta)$$

Likelihood Ratio Policy Gradient

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) = \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau)$$

Estimate this empirically as

$$\approx \hat{g} = (1/m) \sum_{i=1}^m R(\tau^{(i)}) \nabla_{\theta} \log P(\tau^{(i)}; \theta)$$

Remember now this is something like $\mu(s_0)\pi_{\theta}(a_0|s_0)P(s_1|s_0, a_0)\pi_{\theta}(a_1|s_1)\dots$



Likelihood Ratio Policy Gradient

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \sum_{\tau} P(\tau; \theta) R(\tau) = \sum_{\tau} P(\tau; \theta) \nabla_{\theta} \log P(\tau; \theta) R(\tau)$$

Estimate this empirically as

$$\approx \hat{g} = (1/m) \sum_{i=1}^m R(\tau^{(i)}) \nabla_{\theta} \log P(\tau^{(i)}; \theta)$$

Do we need a transition model to compute it?



Score Function

$$\begin{aligned}\nabla_{\theta} \log P(\tau^{(i)}; \theta) &= \nabla_{\theta} \log \left[\underbrace{\mu(s_0)}_{\text{Initial state distrib.}} \prod_{t=0}^{T-1} \underbrace{\pi_{\theta}(a_t|s_t)}_{\text{policy}} \underbrace{P(s_{t+1}|s_t, a_t)}_{\text{dynamics model}} \right] \\ &= \nabla_{\theta} \left[\log \mu(s_0) + \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t|s_t) + \log P(s_{t+1}|s_t, a_t) \right] \\ &= \sum_{t=0}^{T-1} \underbrace{\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)}_{\text{no dynamics model required!}}\end{aligned}$$

NO!

Score Function

$$\begin{aligned}\nabla_{\theta} \log P(\tau^{(i)}; \theta) &= \nabla_{\theta} \log \left[\underbrace{\mu(s_0)}_{\text{Initial state distrib.}} \prod_{t=0}^{T-1} \underbrace{\pi_{\theta}(a_t|s_t)}_{\text{policy}} \underbrace{P(s_{t+1}|s_t, a_t)}_{\text{dynamics model}} \right] \\ &= \nabla_{\theta} \left[\log \mu(s_0) + \sum_{t=0}^{T-1} \log \pi_{\theta}(a_t|s_t) + \log P(s_{t+1}|s_t, a_t) \right] \\ &= \sum_{t=0}^{T-1} \underbrace{\nabla_{\theta} \log \pi_{\theta}(a_t|s_t)}_{\text{no dynamics model required!}}\end{aligned}$$

Called **score function**

Likelihood Ratio Policy Gradient

$$\begin{aligned}\nabla_{\theta} J(\theta) &\approx \hat{g} = (1/m) \sum_{i=1}^m R(\tau^{(i)}) \nabla_{\theta} \log P(\tau^{(i)}; \theta) \\ &= (1/m) \sum_{i=1}^m \boxed{R(\tau^{(i)})} \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t^{(i)} | s_t^{(i)})\end{aligned}$$


Measure how good the sample trajectory is



Visualizing the PG

— — —

$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$$




Visualizing the PG

— — —

$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$$

Moving in the direction the estimated gradient pushes up the logprob of the sample, in proportion to how good it is




Visualizing the PG

— — —

$$\hat{g}_i = f(x_i) \nabla_{\theta} \log p(x_i | \theta)$$

Increase the likelihood of sampling a trajectory with high total reward



Policy Gradient Theorem (Infinite Setting)

The policy gradient theorem generalizes the likelihood ratio approach

$$\begin{aligned}\nabla_{\theta} J(\pi_{\theta}) &= \sum_{h=0}^{\infty} \gamma^h \mathbb{E}_{s, a \sim \mathbb{P}_h^{\pi_{\theta}}} \nabla_{\theta} \ln \pi_{\theta}(a | s) \cdot Q^{\pi_{\theta}}(s, a) \\ &= \frac{1}{1 - \gamma} \mathbb{E}_{s, a \sim d_{\mu}^{\pi_{\theta}}} [\nabla_{\theta} \ln \pi_{\theta}(a | s) \cdot Q^{\pi_{\theta}}(s, a)]\end{aligned}$$



Policy Gradient Theorem (Infinite Setting)

The policy gradient theorem generalizes the likelihood ratio approach

$$\begin{aligned}\nabla_{\theta} J(\pi_{\theta}) &= \sum_{h=0}^{\infty} \gamma^h \mathbb{E}_{s,a \sim \mathbb{P}_h^{\pi_{\theta}}} \nabla_{\theta} \ln \pi_{\theta}(a | s) \cdot Q^{\pi_{\theta}}(s, a) \\ &= \frac{1}{1 - \gamma} \mathbb{E}_{s,a \sim d_{\mu}^{\pi_{\theta}}} [\nabla_{\theta} \ln \pi_{\theta}(a | s) \cdot Q^{\pi_{\theta}}(s, a)]\end{aligned}$$

We want to slowly adjust the policy, such that probability is large at actions with large Q value



Policy Gradient Theorem (Finite Setting)

The policy gradient theorem generalizes the likelihood ratio approach

$$\nabla_{\theta} J(\pi_{\theta}) = \sum_{h=0}^{H-1} \mathbb{E}_{s_h, a_h \sim \mathbb{P}_h^{\pi_{\theta}}} \left[\nabla \ln \pi_{\theta}(a_h | s_h) \cdot Q_h^{\pi_{\theta}}(s_h, a_h) \right]$$

